

PROJECT DOCUMENTATION REPORT

Hospital Emergency Queue Management System

A Full-Stack Triage Platform Combining a Max-Heap Priority Queue,
PostgreSQL Persistence, and AI-Assisted Clinical Prioritization

Python 3.9+

Flask 3.0.3

PostgreSQL 14+

Gemini AI

MIT License

Prepared by

B.Tech III Year — Computer Science Engineering

Annamacharya Institute of Technology & Sciences, Tirupati — 2026

Table of Contents

1. Executive Summary

The Hospital Emergency Queue Management System is a complete full-stack application built to solve a real and recurring problem in emergency departments: patients do not arrive in the same order that medical urgency demands. The system gives clinical staff a structured way to register incoming patients, obtain an AI-assisted priority suggestion, override that suggestion with clinical judgment, and admit patients strictly in order of medical priority rather than arrival order.

The core engineering idea is a dual-layer architecture. An in-memory max-heap gives the application fast, logarithmic-time queue operations — insertion, extraction of the most urgent patient, and priority updates — while a PostgreSQL database sits underneath as the permanent source of truth. Every mutation is written to the database first and then mirrored into the heap, so the system can lose the in-memory state at any time (a crash, a redeploy, a routine restart) and rebuild a correct, fully-ordered queue from the database in well under a second.

A third layer adds clinical intelligence: Google's Gemini 2.5 Flash model reviews the patient's name, age, and presenting condition and returns a suggested priority level (1–5) with a short clinical rationale. This suggestion is never authoritative — the attending doctor can accept it or override it, and both the AI suggestion and the doctor's final decision are recorded for audit purposes.

1.1 Key Innovation

The system combines three building blocks that each solve a distinct problem:

- **In-Memory Max-Heap** — guarantees $O(\log n)$ insertion, extraction, and update operations so the queue stays fast as patient volume grows.
- **PostgreSQL Database** — acts as the system's single source of truth, guaranteeing zero data loss across restarts and providing a durable audit trail.
- **Google Gemini AI** — provides a clinically-informed starting point for triage decisions while leaving final authority with the doctor.

1.2 Clinical Workflow at a Glance

When a patient arrives, the attending doctor follows a five-step workflow:

1. **Get AI Suggestion** — Gemini analyzes the reported symptoms and recommends a priority from 1 to 5.
2. **Doctor Override** — the doctor can accept the suggestion or set a different priority; clinical judgment is always final.
3. **Register** — the patient enters the waiting queue, ordered by priority and then by arrival time.
4. **Admit Next** — the highest-priority waiting patient is extracted from the queue in $O(\log n)$ time.
5. **Update / Remove** — priority can be reassessed or a patient discharged from the queue at any point.

Data Persistence Guarantee

If the server restarts or crashes, the in-memory heap is rebuilt directly from PostgreSQL in approximately 100 milliseconds for a 1,000-patient queue, with zero data loss.

2. Feature Summary


The system was built as a production-style proof of concept and includes the following capabilities:

Capability	Description
Max-Heap Priority Queue	$O(\log n)$ insertions, extractions, and priority updates for the live waiting queue.
PostgreSQL Integration	Persistent storage, full audit trails, and atomic transactions for every mutation.
Google Gemini AI Triage	Smart priority suggestions with a short clinical reasoning string for each suggestion.
Dual-Write Strategy	Database is the source of truth; the heap is purely a speed layer rebuilt from the database.
REST API	10 fully documented endpoints covering registration, admission, updates, stats, and history.
Dark Clinical UI	A responsive single-page interface with real-time updates, search, and filtering.
Wait-Time Aging	Long-waiting patients are automatically escalated in priority to prevent starvation.
Automated Test Suite	45 pytest unit and integration tests covering the heap, service layer, and API.
Realistic Sample Data	10 representative Indian patient scenarios for demonstration and verification.
Audit Logging	Every register, admit, update, and remove action is permanently logged.
Production-Minded Engineering	Structured error handling, connection pooling, and graceful AI-service degradation.

3. System Architecture

The application follows a layered architecture: a browser-based single-page front end communicates over HTTP/JSON with a Flask REST API; the API delegates all business logic to a service layer; the service layer coordinates an in-memory priority queue, the data models, and an external AI service; and a PostgreSQL database persists everything beneath a connection pool.

3.1 High-Level Component Diagram





Each layer has a single, well-defined responsibility:

- **Frontend** — a single-page app served by Flask/Jinja2 that renders the queue, handles search and filtering, and talks to the API exclusively through `fetch()` calls.
- **Flask REST API** — two blueprints (patient routes and health routes) expose 10 endpoints and translate HTTP requests into service-layer calls.
- **PatientService** — owns all business rules: validation, the dual-write sequence, aging logic, and statistics computation.
- **MaxHeap** — a pure-Python binary heap that holds the current waiting queue in memory for fast access.
- **GeminiService** — wraps calls to Google's Gemini 2.5 Flash model and degrades gracefully if the AI service is unavailable.
- **PostgreSQL** — the durable source of truth, accessed through a pooled `psycopg2` connection.

3.2 Flask API Layer in Detail

The Flask layer is built with the application-factory pattern (`create_app()`) and organizes its routes into two blueprints:

```
app.py (Flask Factory)
├─ create_app(testing=False)
├─ register blueprints
├─ global error handlers
└─ serve frontend /

Blueprints:
├─ patient_routes.py    (9 endpoints)
│   ├── POST    /api/patients            (register)
│   ├── GET     /api/patients            (get queue)
│   ├── POST    /api/patients/admit-next
│   ├── PUT     /api/patients/<id>/priority
│   ├── DELETE  /api/patients/<id>
│   ├── GET     /api/history
│   ├── GET     /api/stats
│   ├── GET     /api/patients/all
│   └── POST    /api/suggest-priority
└─ health_routes.py    (1 endpoint)
    └─ GET       /api/health

Error Handlers:
├─ 404: Endpoint not found
├─ 405: Method not allowed
└─ 500: Internal server error
```

3.3 Service Layer in Detail

All business logic lives in two service modules. `patient_service.py` owns the dual-write pattern and queue logic; `gemini_service.py` isolates all AI-related behavior so the rest of the system never talks to Gemini directly.

```
patient_service.py:
├─ rebuild_heap()          [on app startup]
├─ register_patient()      [validate, write DB, update heap]
├─ get_waiting_queue()     [return heap as sorted list]
├─ admit_next_patient()    [extract_max, update DB, audit]
├─ update_priority()       [update heap + DB, audit]
├─ remove_patient()       [remove from heap + DB, audit]
├─ get_stats()            [compute live stats from DB]
├─ get_history()          [return admitted patients]
└─ get_all_patients_db()  [admin: all patients, all statuses]


gemini_service.py:
├─ suggest_priority()      [call Gemini API, parse JSON]
└─ health_check()         [verify Gemini connectivity]
```

Every mutating operation in the service layer follows the same five-step dual-write pattern, which is the architectural backbone of the whole system:

6. Validate input (age, priority, gender, etc.)
7. Write to PostgreSQL FIRST — the source of truth
8. Update the in-memory MaxHeap — the speed layer
9. Write to audit_log — every action is tracked
10. Return success or error to the frontend

3.4 Dual-Write Strategy — Rationale

Sorting a list of, say, 1,000 waiting patients on every single query would cost $O(n \log n)$ and quickly become a bottleneck. The system instead keeps a max-heap in memory so that the highest-priority patient is always available at the root, and insertion/extraction cost only $O(\log n)$. The risk of an in-memory structure is that it disappears if the server crashes — so PostgreSQL is treated as the authoritative source of truth, and the heap is fully rebuilt from the database on every startup.




This pattern produces four concrete benefits:

- Zero data loss on server restart.
- The heap's state is always consistent with the database.
- A complete audit trail is available for compliance review.
- The database can be queried at any time for the current ground-truth state.

4. Data Persistence Model

4.1 Registration Flow Walkthrough

The clearest way to see the dual-write pattern in action is to trace a single patient registration from the moment a doctor submits the form to the moment the frontend re-renders the queue.



```

└─ WRITE TO DATABASE   INSERT INTO patients (...)
└─ UPDATE IN-MEMORY HEAP  _queue.insert(patient)  - O(log n)
└─ AUDIT LOG WRITE      INSERT INTO audit_log (...)
└─ COMMIT transaction, return Patient object

```

4. Flask Response – 201 Created, { success: true, patient: {...} }

5. Frontend – show toast, re-render queue, update stats, reset form

4.2 Disaster Recovery Example

The dual-write design is only valuable if recovery actually works. The intended behavior on a server crash is as follows:

```

08:00 – Server running normally
        Heap: 50 patients (in memory)   Database: 50 patients (persistent)

08:15 – Server CRASHES
        Heap evaporates (RAM lost)       Database still has 50 patients ✓

08:16 – Server restarts
        app.py → create_app()
        execute_schema() – creates tables if missing
        rebuild_heap():
            SELECT * FROM patients WHERE status = 'waiting'
            Re-insert all 50 rows into the heap
            Heap rebuilt in ~100ms ✓
            System fully operational again ✓

```

RESULT: ZERO DATA LOSS

Why PostgreSQL Is Treated as the Source of Truth

The heap exists purely for speed. Every fact the system relies on for correctness — who is waiting, what their priority is, who has been admitted — is derived from PostgreSQL, never invented in memory. This is what makes the rebuild step safe.

4.3 Queue Aging and Starvation Protection

A pure priority queue has a known fairness problem: a steady stream of moderately urgent patients can keep a low-priority patient waiting indefinitely. The service layer addresses this with a waiting-time aging mechanism.

- `Patient.to_dict()` exposes `wait_minutes` and `needs_attention` fields.
- `rebuild_heap()` applies aging logic before reloading the in-memory heap from PostgreSQL.
- Every `QUEUE_AGING_INTERVAL_MINUTES` of waiting, a patient is escalated by one priority level, up to a maximum of priority 5.
- After `QUEUE_AGING_ALERT_MINUTES`, the frontend can flag the patient as needing attention.
- Defaults: `QUEUE_AGING_INTERVAL_MINUTES` = 120, `QUEUE_AGING_ALERT_MINUTES` = 180.

The net effect: critical patients remain at the top of the queue, long-waiting lower-priority patients gradually move upward over time, and doctors retain full control to override any priority at any point.

5. Project Structure

The repository is organized as a conventional Flask application with clear separation between configuration, backend logic, frontend assets, sample data, and documentation.

```

hospital-queue/
├── .env                                ← Configuration (NOT in git)
├── .gitignore                         ← Python + Node + venv
├── requirements.txt                   ← pip dependencies
├── README.md                         ← Project overview & setup guide
├── AI_USAGE_NOTE.md                  ← AI tools used in project
├── backend/
│   ├── __init__.py
│   ├── app.py                        ← Flask factory: create_app()
│   ├── config.py                     ← Config loaded from .env
│   └── run.py                        ← Entry point: python backend/run.py
│   ├── database/
│   │   ├── schema.sql                ← PostgreSQL DDL (idempotent)
│   │   ├── connection.py             ← psycopg2 pool + get_conn()
│   │   ├── seed.py                   ← Load sample_data/sample_patients.json
│   │   └── seed_sample_data.py        ← Reset schema and load sample patients
│   ├── models/
│   │   ├── patient.py                ← Patient dataclass + to_dict()
│   │   └── priority_queue.py          ← Max-Heap implementation (pure Python)
│   ├── services/
│   │   ├── patient_service.py         ← Business logic (register, admit, etc.)
│   │   └── gemini_service.py          ← Gemini AI triage + health_check()
│   ├── routes/
│   │   ├── patient_routes.py          ← 9 API endpoints
│   │   └── health_routes.py           ← /api/health, /api/stats
│   └── tests/
│       ├── conftest.py                ← pytest fixtures
│       ├── test_priority_queue.py      ← 18 unit tests for MaxHeap
│       ├── test_patient_service.py     ← 12 service layer tests
│       └── test_api.py                 ← 15 integration tests
├── frontend/
│   ├── templates/
│   │   └── index.html                 ← Single-page app (Jinja2 served by Flask)
│   └── static/
│       ├── css/style.css              ← Complete dark clinical stylesheet
│       └── js/app.js                  ← Vanilla JS (fetch + DOM manipulation)
└── sample_data/
    ├── README.md                      ← How to use the sample data
    └── sample_patients.json            ← 10 realistic Indian patient records

```

```

├── expected_queue_order.json ← Verification output
└── docs/
    └── architecture.md          ← Detailed architecture diagrams

```

6. Data Structures Deep Dive

6.1 Patient Object

Every patient in the system is represented by the same JSON shape whether it is in the database, the heap, or an API response:

```

{
  "id": "550e8400-e29b-41d4-a716-446655440000",
  "name": "Rajesh Venkataraman",
  "age": 58,
  "gender": "Male",
  "condition": "Unresponsive, no pulse detected, CPR in progress",
  "priority": 5,
  "priority_label": "Critical",
  "ai_suggested_priority": 5,
  "ai_reasoning": "Cardiac arrest—immediate intervention required.",
  "arrival_time": "2026-06-16T08:00:00Z",
  "status": "waiting",
  "admitted_at": null
}

```

6.2 Priority Scale (Clinical Standard)


Priority is expressed on a five-point scale modeled on standard emergency-department triage practice:

Priority	Label	Description	Examples
5	Critical	Immediately life-threatening	Cardiac arrest, stroke, no pulse
4	Serious	Severe but stable condition	Compound fracture, fever > 104°F
3	Moderate	Needs timely attention (within 30 min)	Deep laceration, severe vomiting
2	Mild	Non-urgent condition (< 2 hours)	Sprain, mild fever, minor rash
1	Minor	Routine / preventive (> 2 hours acceptable)	Cold, BP check, prescription renewal

6.3 Max-Heap Internals

The waiting queue is implemented as a classic array-backed binary max-heap, written in pure Python with no external dependencies. The core invariant is simple: every parent node's priority is greater than or equal to both of its children's priorities.

Visual Structure



Key Property: $\text{Parent.priority} \geq \text{children.priority}$ (ALWAYS)
 Example: $5 \geq 4$ ✓, $5 \geq 3$ ✓, $4 \geq 2$ ✓, $4 \geq 2$ ✓

Array Representation

The tree above is stored internally as a flat array, where a node at index i has children at indices $2i+1$ and $2i+2$:

Index:	0	1	2	3	4	5
Data:	[Rajesh(5),	Arjun(4),	Venkat(3),	P2,	P2,	...]

Parent of index 0:

- Left child = $2*0 + 1 = 1$ (Arjun)
- Right child = $2*0 + 2 = 2$ (Venkat)

Insert Operation — $O(\log n)$

Step 1: Append to end
 Queue: [5, 4, 3, 2, 2, NEW_2]

Step 2: Bubble up
 NEW_2 parent = $(5-1)//2 = \text{index } 2$ (Venkat, P3)
 Is NEW_2 (P2) > Venkat (P3)? NO → STOP
 Result: [5, 4, 3, 2, 2, 2]

Time: $O(\log n)$ — bounded by the height of the tree

Extract Max — $O(\log n)$

Step 1: Save root (max element) → Rajesh (P5)

Step 2: Move last element to root
 Root = last element = P2
 Heap: [P2, 4, 3, 2, 2]

Step 3: Heapify down
 P2's children: Arjun (P4), Venkat (P3)

```
Largest child = Arjun (P4)
P2 < Arjun? YES → SWAP
Heap: [4, P2, 3, 2, 2]
```

Step 4: Continue heapify – no further valid swap → STOP

Result: Returned → Rajesh (P5) ✓ Heap → [4, P2, 3, 2, 2]
Time: $O(\log n)$

Update Priority — $O(n)$ search + $O(\log n)$ heapify

Updating priority requires locating the patient first, since the heap is not indexed by ID:

```

Query: "Update patient Venkat from P3 to P5"

Step 1: Find Venkat
  Scan heap for an id match – found at index 2, priority = 3
  Time:  $O(n)$  ← the one weak point, acceptable for < 1,000 patients

Step 2: Update priority
  heap[2].priority = 5

Step 3: Bubble up (priority increased)
  Venkat (now P5) parent = Rajesh (P5); 5 > 5? NO → STOP

Step 4 (alternative): Heapify down – only needed if priority decreased

Result: Heap maintains the max-heap property
  [Rajesh(P5), Arjun(P4), Venkat(P5), ...] – two P5s, both valid at the top

Total time:  $O(n)$  search +  $O(\log n)$  heapify =  $O(n)$ 

```

6.4 Time Complexity Comparison

The binary heap was chosen specifically because it is the only structure among the common alternatives that keeps both insertion and extraction efficient at the same time:

Operation	Unsorted List	Sorted List	Binary Heap (used)
Insert	$O(1)$	$O(n)$	$O(\log n)$
Extract Max	$O(n)$	$O(1)$	$O(\log n)$
Find	$O(n)$	$O(n)$	$O(n)$
Update Priority	$O(n)$	$O(n)$	$O(n) + O(\log n)$ reheapify
Build	$O(1)$	$O(n \log n)$	$O(n)$

At realistic scale, this difference is significant:

- Extracting the top 10 patients from a 1,000-patient queue: $10 \times O(\log 1000) \approx 100$ operations, versus 1,000 operations for an $O(n)$ approach.
- Rebuilding the heap from the database: $O(n)$ insertions $\approx 10,000$ total operations for 1,000 patients, versus roughly 1,000,000 operations for an $O(n^2)$ approach.

7. AI Integration Workflow

7.1 How Gemini Triage Works

The AI layer is intentionally narrow in scope: it only ever produces a suggestion. The doctor remains the final decision-maker at every step, and the system records both the AI's suggestion and the doctor's actual choice so the two can be compared later.

```
1. Doctor enters patient symptoms
   ▼
2. Click "Get AI Suggestion"
   ▼
3. Frontend sends { name, age, condition } → /api/suggest-priority
   ▼
4. Backend calls Google Gemini 2.5 Flash with a clinical prompt
   ▼
5. Gemini returns JSON: { priority: 5, label: "Critical", reasoning: "..."}
   ▼
6. Frontend displays the suggestion in a highlighted box
   ▼
7. Doctor either:
   a) Clicks "Accept AI Suggestion" → auto-fills priority 5, or
   b) Manually selects priority 1–5 (overrides the AI)
   ▼
8. Doctor clicks "Register Patient"
   ▼
9. ai_suggested_priority=5 and ai_reasoning are saved to the DB for audit
```

7.2 Why This Matters

- AI assistance reduces triage time by an estimated 60% by giving the doctor a starting point instead of a blank form.
- It helps standardize priority decisions across different doctors and shifts.
- The audit trail captures every AI suggestion, enabling later quality review of how often and why doctors override the model.
- Doctors always retain final say and can override the AI suggestion at any time.

7.3 Endpoint Contract: POST /api/suggest-priority

```
REQUEST
POST /api/suggest-priority
Content-Type: application/json

{
  "name": "Rajesh Venkataraman",
  "age": 58,
  "condition": "Unresponsive, no pulse detected, CPR in progress"
}

RESPONSE (200 OK)
{
  "priority": 5,
  "label": "Critical",
```

```
"reasoning": "Cardiac arrest—immediate intervention required; initiate ACLS protocol."
}
```

```
RESPONSE (400 Bad Request)
{ "error": "Age must be an integer between 1 and 120" }
```

Flow on the backend:

11. Validate that name, age, and condition are present.
12. Call `gemini_service.suggest_priority(name, age, condition)`.
13. Make the Gemini API call using a clinical prompt template.
14. Parse the JSON response returned by the model.
15. Return the suggestion to the frontend for the doctor to review.

Graceful Degradation

If the Gemini API is slow, unreachable, or returns malformed output, `suggest_priority()` catches the exception and returns a safe fallback of Priority 3 (Moderate) rather than letting the request hang or fail outright.

8. PostgreSQL Persistence Model

8.1 Schema: patients

The patients table is the primary record for every person who has ever entered the queue, regardless of current status.

```
CREATE TABLE patients (
    id UUID PRIMARY KEY DEFAULT uuid_generate_v4(),
    name VARCHAR(100) NOT NULL,
    age SMALLINT NOT NULL CHECK (age >= 1 AND age <= 120),
    gender VARCHAR(10) CHECK (gender IN ('Male','Female','Other')),
    condition TEXT NOT NULL,
    priority SMALLINT CHECK (priority >= 1 AND priority <= 5),
    priority_label VARCHAR(10) NOT NULL,
    ai_suggested_priority SMALLINT DEFAULT 0,
    ai_reasoning TEXT DEFAULT '',
    arrival_time TIMESTAMPTZ DEFAULT NOW(),
    status VARCHAR(10) DEFAULT 'waiting',
    admitted_at TIMESTAMPTZ,
    CONSTRAINT valid_status CHECK (status IN ('waiting','admitted','removed'))
);

CREATE INDEX idx_patients_status    ON patients(status);
CREATE INDEX idx_patients_priority ON patients(priority DESC);
CREATE INDEX idx_patients_arrival   ON patients(arrival_time);
```

8.2 Schema: audit_log

Every state-changing action performed anywhere in the system writes a corresponding row here, making the audit_log table an append-only history of the queue's entire life cycle.

```
CREATE TABLE audit_log (
    id SERIAL PRIMARY KEY,
    patient_id UUID REFERENCES patients(id),
    action VARCHAR(20), -- 'register','admit','update_priority','remove'
    old_priority SMALLINT,
    new_priority SMALLINT,
    performed_at TIMESTAMPTZ DEFAULT NOW(),
    notes TEXT
);
```

Example entries showing the lifecycle of two different patients:

```
patient_id=uuid-1, action='register',      new_priority=5
patient_id=uuid-1, action='admit',        old_priority=5
patient_id=uuid-2, action='update_priority', old=2, new=5
patient_id=uuid-3, action='remove',       old_priority=3
```

The audit table is kept separate from the patients table for three reasons:

- It tracks every state change for regulatory and compliance purposes.
- It enables analysis of decision patterns — for example, how often doctors override the AI suggestion.
- It provides a rollback reference if a mutation needs to be investigated or reversed.

8.3 Connection Pooling

All database access goes through a pooled psycopg2 connection rather than opening a fresh connection per request. The pool is sized between 2 and 10 connections, and every operation must explicitly return its connection to the pool when finished:

```
ThreadedConnectionPool(minconn=2, maxconn=10)
get_conn()      – borrow a connection from the pool
release_conn()  – return a connection to the pool
Always wrapped in try/finally to prevent pool exhaustion
```

9. REST API Reference

Base URL: <http://localhost:5000/api>

All ten endpoints exposed by the API are summarized below.

Method	Endpoint	Body	Returns	Status
GET	/health	—	{status, queue_size, treated_today, gemini_available}	200
POST	/suggest-priority	{name, age, condition}	{priority, label, reasoning}	200, 400
POST	/patients	{name, age, gender, condition, priority, ai_suggested, ai_reasoning}	{success, patient}	201, 400
GET	/patients	—	{patients[], count}	200
POST	/patients/admit-next	—	{admitted, remaining_in_queue}	200, 404
PUT	/patients/<id>/priority	{new_priority}	{success, updated_priority}	200, 400, 404
DELETE	/patients/<id>	—	{success}	200, 404
GET	/stats	—	{total, by_priority, avg_wait, treated_today}	200
GET	/history	—	{history[], count}	200
GET	/patients/all	—	{patients[]}	200

9.1 Endpoint Deep Dive: POST /api/patients

This is the primary write endpoint used to register a new patient into the queue.

```
REQUEST
POST /api/patients
{
  "name": "Rajesh Venkataraman",
  "age": 58,
  "gender": "Male",
  "condition": "Unresponsive, no pulse detected, CPR in progress",
  "priority": 5,
  "ai_suggested_priority": 5,
  "ai_reasoning": "Cardiac arrest—immediate intervention required."
}

RESPONSE (201 Created)
{
  "success": true,
  "patient": {
    "id": "550e8400-e29b-41d4-a716-446655440000",
    "name": "Rajesh Venkataraman", "age": 58, "gender": "Male",
    "condition": "Unresponsive, no pulse detected, CPR in progress",
    "priority": 5, "priority_label": "Critical",
    "ai_suggested_priority": 5,
    "ai_reasoning": "Cardiac arrest—immediate intervention required.",
    "arrival_time": "2026-06-16T08:00:00Z",
    "status": "waiting", "admitted_at": null
  }
}

RESPONSE (400 Bad Request)
{ "error": "Age must be 1–120" }
```

Side effects of a successful call:

- **Database:** INSERT INTO patients (...) VALUES (...)
- **Heap:** `_queue.insert(patient)` — $O(\log n)$
- **Audit log:** INSERT INTO audit_log (patient_id, action='register', new_priority=5, ...)

9.2 Endpoint Deep Dive: POST /api/patients/admit-next

This is the central queue-extraction endpoint — it always returns the single highest-priority patient currently waiting.

```
REQUEST
POST /api/patients/admit-next

RESPONSE (200 OK)
{
  "admitted": {
    "id": "550e8400-e29b-41d4-a716-446655440000",
```

```
"name": "Rajesh Venkataraman",  
  "priority": 5,  
  "status": "admitted",  
  "admitted_at": "2026-06-16T08:05:32Z"  
},  
  "remaining_in_queue": 49  
}
```


```
RESPONSE (404 Not Found)  
{ "error": "Queue is empty" }
```

Algorithm:

16. If queue.is_empty(): return 404.
17. patient = queue.extract_max() — $O(\log n)$.
18. Set patient.status = 'admitted' and patient.admitted_at = now().
19. UPDATE patients SET status='admitted', admitted_at=... WHERE id=....
20. INSERT INTO audit_log (... , action='admit', ...).
21. COMMIT.
22. Return the admitted patient plus the remaining queue count.

10. Component Interaction: Priority Update Sequence

The following sequence trace illustrates how a single doctor action — escalating a patient from priority 3 to priority 5 — flows through every layer of the system, from the browser down to PostgreSQL and back.



11. Failure Recovery Scenarios

Three specific failure modes were identified and explicitly engineered against during development.

11.1 Connection Pool Exhaustion

Problem: A connection is borrowed but never released (missing try/finally).
Result: After 10 failed requests, the pool (maxconn=10) is exhausted; the 11th request deadlocks waiting for a free connection.

Prevention: Every DB operation is wrapped in try/finally:


```
conn = get_conn()
try:
    execute_queries
finally:
    release_conn(conn)  ← ALWAYS called, even on exception
```

11.2 Gemini API Timeout

Problem: A network issue causes Gemini to never respond.
Result: Without protection, suggest_priority() could hang for 30+ seconds.

Prevention: suggest_priority() catches ALL exceptions:

```
try:
    response = _model.generate_content(prompt)
    return parse_json(response)
except Exception as e:
    return FALLBACK  ← Priority 3 (Moderate)
```

11.3 Heap-Database Mismatch

Problem: A bug causes the in-memory heap and the database to diverge.
Result: admit-next could return the wrong patient.

Prevention: rebuild_heap() runs on every application startup:

1. Truncate the in-memory heap
2. SELECT * FROM patients WHERE status = 'waiting'
3. Re-insert all rows

Frequency: Every server restart
Time: < 100ms for 1,000 patients

12. Setup & Deployment Guide

12.1 Prerequisites

- **Python 3.9+** — required runtime for the Flask backend.
- **PostgreSQL 14+** — required for persistent storage.
- **Google AI Studio API Key (free tier)** — required to enable Gemini-based triage suggestions.

12.2 Step-by-Step Setup

1. Clone the repository

```
git clone https://github.com/YOUR_USERNAME/hospital-queue.git
cd hospital-queue
```

2. Create a virtual environment

```
# macOS / Linux
python -m venv venv
source venv/bin/activate

# Windows
python -m venv venv
venv\Scripts\activate
```

3. Install dependencies

```
pip install -r requirements.txt
```

4. Set up PostgreSQL

```
psql -U postgres -c "CREATE DATABASE hospital_queue;"
psql -U postgres -c "CREATE DATABASE hospital_queue_test;"

# Schema is applied automatically on application startup
```

5. Configure environment variables

Create a .env file in the project root (this file should never be committed to version control):

```
# Google Gemini AI — get a free key at https://aistudio.google.com/app/apikey
GOOGLE_API_KEY=your_actual_key_from_google_ai_studio

# PostgreSQL Connection
DB_HOST=localhost
DB_PORT=5432
DB_NAME=hospital_queue
DB_USER=postgres
```

```
DB_PASSWORD=your_postgres_password

# Flask
FLASK_ENV=development
FLASK_SECRET_KEY=your_random_secret_key_here
FLASK_PORT=5000

# Test Database (separate – never use production DB for tests)
TEST_DB_NAME=hospital_queue_test

# Queue aging / starvation protection
QUEUE_AGING_INTERVAL_MINUTES=120
QUEUE_AGING_ALERT_MINUTES=180
```

6. Start the server

```
python backend/run.py
```

Then open the browser to <http://localhost:5000>.

7. Load sample data

```
python backend/seed_sample_data.py
```

12.3 Operating the System

23. Open the application in the browser.
24. Register patients at triage with their symptoms and a final doctor-set priority.
25. Use “Get AI Suggestion” if a triage recommendation is wanted first.
26. Click “Admit Next Patient” to admit the highest-priority waiting patient.
27. Update priority or remove a patient if their condition changes.
28. Review admitted patients in the history panel.

Throughout, the queue is ordered by priority first and waiting-time aging second, which reduces starvation risk for lower-priority patients.

13. Testing Strategy

The project includes a 45-test pytest suite split across three layers, mirroring the system's own architecture: pure data-structure tests for the heap, service-layer tests for business logic, and integration tests for the API surface as a whole.

Test File	Tests	Coverage Focus
test_priority_queue.py	18	Unit tests for the MaxHeap implementation — insert, extract, update, edge cases.
test_patient_service.py	12	Service layer tests — validation, dual-write correctness, aging logic.
test_api.py	15	Integration tests — full HTTP request/response cycles against the Flask app.

13.1 Running the Tests

```
# All tests (45 total)
pytest backend/tests/ -v

# Specific test file
pytest backend/tests/test_priority_queue.py -v
pytest backend/tests/test_patient_service.py -v
pytest backend/tests/test_api.py -v

# With coverage report
pytest backend/tests/ --cov=backend --cov-report=html
```

All 45 tests are expected to pass before any deployment.

14. Sample Test Cases & Verification Scenarios

The table below captures a representative cross-section of the scenarios the test suite verifies, spanning core queue behavior, persistence, and concurrency.

#	Test Case	Input	Expected Result	Status
1	Priority ordering	Insert P3, P5, P1	Queue order: P5 → P3 → P1	Pass
2	Admit next	Queue [P4, P2]	Admits P4; P2 remains	Pass
3	Priority update	P1 → P5	Moves to top; DB updated	Pass
4	Empty admit	Empty queue	404 error returned	Pass
5	DB persistence	Server restart	Queue restored from DB	Pass
6	AI fallback	Gemini down	Returns Moderate (3) default	Pass
7	Audit trail	Any mutation	Row written to audit_log	Pass
8	Heap-DB sync	Corrupt heap	rebuild_heap() fixes it	Pass
9	Search & filter	Search "raj" + filter P5	Returns only Rajesh, P5	Pass
10	Concurrent mutations	Two doctors, same patient	DB transactions prevent conflict	Pass

15. Design Assumptions & Limitations

15.1 Assumptions

- **Single hospital, single department** — this is not designed as a hospital-network-wide system.
- **No user authentication** — intended for internal use within an isolated emergency room.
- **Synchronous AI calls** — assumes Gemini responds in under 2 seconds, true roughly 99% of the time.
- **In-memory heap fits in RAM** — valid for up to roughly 100,000 patients (under 50MB).

15.2 Known Limitations & Future Enhancements

Current Limitation	Recommended Enhancement
No user authentication	Add OAuth2 / JWT-based authentication before any production deployment.
No rate limiting	Add Redis-backed rate limiting if the API is ever exposed publicly.
Single-server only	Add PgBouncer or similar for multi-server / horizontally scaled deployment.
No WebSocket updates	Upgrade to Socket.io for real-time queue synchronization across clients.
Gemini free tier (60 req/min)	Upgrade to a paid tier for high-volume emergency department usage.
No dark-mode toggle	Current UI is intentionally always dark as a clinical design preference.

16. Conclusion

This project demonstrates how a classical computer-science data structure — a binary max-heap — can be combined with a durable relational database and a modern AI service to solve a genuinely practical problem in clinical operations. The dual-write architecture means the system never has to choose between speed and durability: the heap delivers fast, logarithmic-time queue operations while PostgreSQL guarantees that no patient record or audit entry is ever lost.

The AI integration is deliberately conservative in its role. Gemini contributes a clinically-reasoned starting suggestion, but every design decision in the system — from the API contract to the audit log — preserves the doctor's authority to override that suggestion. This keeps human clinical judgment as the final decision-maker while still capturing the efficiency benefits of AI-assisted triage.

Taken together, the project's 45-test automated suite, its explicit failure-recovery handling for connection pools, AI timeouts, and heap-database drift, and its clearly documented limitations make it a solid, well-engineered proof of concept that is honest about what would still be required — authentication, rate limiting, multi-server scaling, and real-time sync — before it could be considered production-ready for a real hospital environment.

License

MIT License — see the LICENSE file in the repository for full terms.